

The Missing Benchmark Metric: Memory

`#java #benchmark #JMH #memory #metrics #GC
#caching #optimizethewhole #bigpicture`

2nd February 2019
Java Dev Room, FOSDEM, Brussels

Jens Wilke
@cruftex
cruftex.net

CC BY-SA 3.0

About Me



<https://cache2k.org>

- Jens Wilke
- Performance Fan(atic)
- Author of cache2k
- 70+ answered questions on StackOverflow about Caching
- JCache / JSR107 Contributor
- Tiny contributions to JMH
- @cruftex / cruftex.net

Content

- Why?
- How to gather memory metrics?
- How (not) to gather memory metrics from a running JMH benchmark?
- Real benchmark results ...
.... to get a feeling whether this is worthwhile
- Happyness?!
- Please ask questions right away!

Motivation

- Compare different caching libraries in identical scenario (access sequence, cache size (entries!), CPU and JDK) with JMH
- Primary result: Throughput in **ops/s**
- But: Compare **caching libraries** and ignore memory consumption?!
- Space vs. Time trade off
- But:
no mechanism in JMH to record memory usage metrics (yet...)

Metric: Object Graph Traversing

- EHCafe sizeof
- Java Agent for Memory Measurements

Pro:

- No actual usage long running benchmark needed, just fill up the data structures
- No complex setup (dedicated hardware)

Con:

- How to choose the root object / traverse the whole heap?
- Static result
- heap objects only
- Needs time

Metric: Heap Dump / Heap Histogram

- e.g. use jmap to get a heap dump or histogram


Pro:

- Same as before, we don't need a real benchmark run
- Accurate value of used heap space
- Additional information in the histogram or dump


Con:

- Costly
- Static value of heap objects only


Let it Run!




How to Extract a (single) Metric?



Fore GC?




Side to Side Comparison




^ no forced GC

forced GC >



Wait a minute....



Sidenote:


GC and Micro Benchmarks

- Micro-Microbenchmark: GC might happen during an iteration
 - Single GC occurrence causes skewed result
 - Use JMH parameter `-gc true` or Zero GC


=> Get garbage collector out of the equation
- NotSo-Microbenchmark: GC happens always during an iteration
 - Make sure a lot of GCs happen, JMH parameter: `-prof gc`
 - Increase warmup and iteration time
 - „Know your GC“ \leftrightarrow „Get to know your GC“

=> You cannot avoid the GC. Make it go steady.

A Metric? Some Metrics!




Linux OS Metrics: VmRSS, VmHWM




GC Notifications

- Notification after GC run (via `GarbageCollectionNotificationInfo`)
- Contains memory usage before and after GC
(examples only shows notification in iteration 2)




Allocation Rate



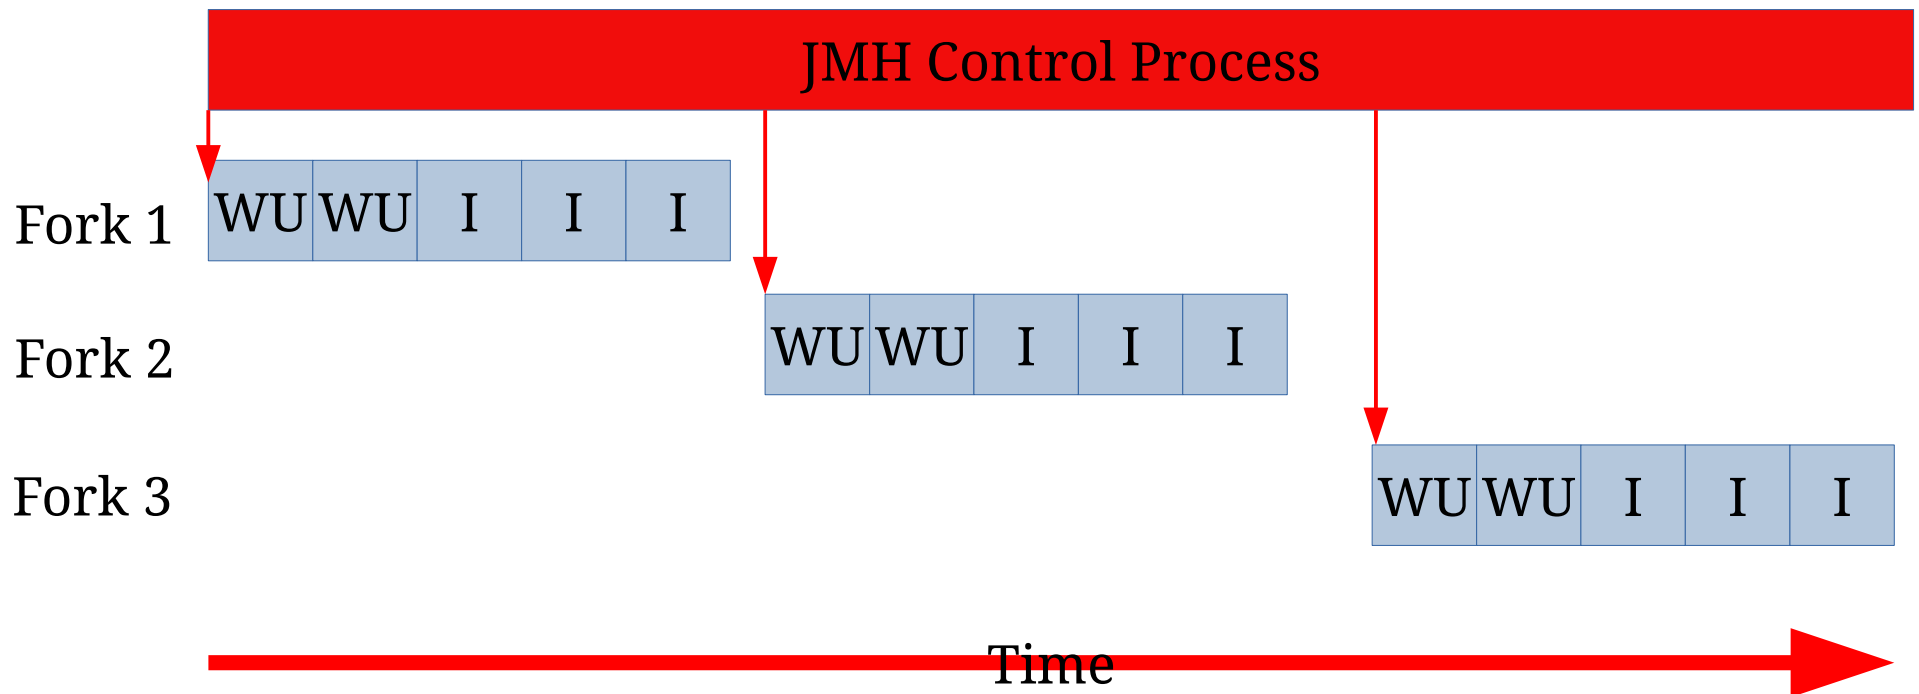
Minimal Memory at the End

- JMH can run multiple VMs after another (forks)
- at end of each fork the invasive collectors could happen (force GC, heap histogram)



Example JMH Running Scheme

- JMH can run multiple VMs after another (forks)
- 9 iterations to collect results
- 3 forks, at end of each fork the invasive collectors happen (force GC, heap histogram)
- Parameters: `-f 3 -wi 3 -i 3`



Results

JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;

@State(Scope.Thread)
public static class ThreadState {

    ZipfianPattern pattern;

    @Setup(Level.Iteration)
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }

}

@Benchmark @BenchmarkMode(Mode.Throughput)
public long operation(ThreadState threadState, HitCountRecorder rec) {
    // TODO: JMH should return the raw number of operations somewhere...
    rec.opCount++;
    Integer v = cache.get(threadState.pattern.next());
    return v;
}

public Integer load(final Integer key) {
    missCount.increment();
    Blackhole.consumeCPU(1000);
    return key * 2 + 11;
}
```

JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;
```

```
@State(Scope.Thread)
```

```
public static class ThreadState {
```

```
    ZipfianPattern pattern;
```

```
    @Setup(Level.Iteration)
```

```
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }
```

```
}
```

```
@Benchmark @BenchmarkMode(Mode.Throughput)
```

```
public long operation(ThreadState threadState, HitCountRecorder rec) {
```

```
    // TODO: JMH should return the raw number of operations somewhere...
```

```
    rec.opCount++;
```

```
    Integer v = cache.get(threadState.pattern.next());
```

```
    return v;
```

```
}
```

```
public Integer load(final Integer key) {
```

```
    missCount.increment();
```

```
    Blackhole.consumeCPU(1000);
```

```
    return key * 2 + 11;
```

```
}
```

(fast) Zipfian sequence generator per thread
Skewed random pattern, yields around 90% hitrate

JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;
```

```
@State(Scope.Thread)
```

```
public static class ThreadState {
```

```
    ZipfianPattern pattern;
```

```
    @Setup(Level.Iteration)
```

```
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }
```

```
}
```

```
@Benchmark @BenchmarkMode({Mode.Throughput})
```

```
public long operation(ThreadState threadState, HitCountRecorder rec) {
```

```
    // TODO: JMH should return the raw number of operations somewhere...
```

```
    rec.opCount++;
```

```
    Integer v = cache.get(threadState.pattern.next());
```

```
    return v;
```

```
}
```

```
public Integer load(final Integer key) {
```

```
    missCount.increment();
```

```
    Blackhole.consumeCPU(1000);
```

```
    return key * 2 + 11;
```

```
}
```

(fast) Zipfian sequence generator per thread
Skewed random pattern, yields around 90% hitrate

Benchmarked operation is cache.get()
On miss, cache will invoke a load function(read-through)

JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;
```

```
@State(Scope.Thread)
public static class ThreadState {
```

```
    ZipfianPattern pattern;
```

```
    @Setup(Level.Iteration)
```

```
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }
```

```
}
```

```
@Benchmark @BenchmarkMode({Mode.Throughput})
```

```
public long operation(ThreadState threadState, HitCountRecorder rec) {
    // TODO: JMH should return the raw number of operations somewhere...
    rec.opCount++;
    Integer v = cache.get(threadState.pattern.next());
    return v;
}
```

```
public Integer load(final Integer key) {
    missCount.increment();
    Blackhole.consumeCPU(1000);
    return key * 2 + 11;
}
```

(fast) Zipfian sequence generator per thread
Skewed random pattern, yields around 90% hitrate

Benchmarked operation is cache.get()
On miss, cache will invoke a load function(read-trough)

Load function has penalty, via
Blackhole.consumeCPU()

JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;
```

```
@State(Scope.Thread)
```

```
public static class ThreadState {
```

```
    ZipfianPattern pattern;
```

```
    @Setup(Level.Iteration)
```

```
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }
```

```
}
```

```
@Benchmark @BenchmarkMode({Mode.Throughput})
```

```
public long operation(ThreadState threadState, HitCountRecorder rec) {
```

```
    // TODO: JMH should return the raw number of operations somewhere...
```

```
    rec.opCount++;
```

```
    Integer v = cache.get(threadState.pattern.next());
```

```
    return v;
```

```
}
```

```
public Integer load(final Integer key) {
```

```
    missCount.increment();
```

```
    Blackhole.consumeCPU(1000);
```

```
    return key * 2 + 11;
```

```
}
```

(fast) Zipfian sequence generator per thread
Skewed random pattern, yields around 90% hitrate

Benchmarked operation is cache.get()
On miss, cache will invoke a load function(read-trough)

Load function has penalty, via
Blackhole.consumeCPU()

=> a lot can happens on cache.get()
=> Cache eviction and autoboxing

Benchmark Setup

Environment:

- CPU: Intel(R) Xeon(R) CPU E3-1240 v5 @ 3.50GHz, **4 physical cores**
 - Benchmarks run with 4 thread, 4 cores (limited via CPU hotplug)
- Oracle JDK 11, ParallelGC and G1
- Ubuntu 18.04

Cache Library Versions:

- Google Guava Cache, Version 26
- Caffeine, Version 2.6.2
- cache2k, Version 1.2.0.Final
- EHCache, Version 3.6.1


Code is at:

<https://github.com/cache2k/cache2k-benchmark>


JMH Parameters

- JMH Parameters:
 - 3 forks, 2 warmup iterations, 3 measurement iterations, 60 seconds iterations time
- => 9 measurement iterations
- Graphs show the confidence interval
- Confidence interval is at 99.9% confidence level!


Parallel GC: ops/s




Parallel GC: heap usage




Parallel GC: Total / RSS




G1: heap usage




G1: Total / RSS



G1: allocation rate op/s



G1: allocation rate byte/op



Conclusion

- If GC is happening, also record memory usage in your JMH results
- JVM / OS memory metrics can be utilized, with varying degree of truth and consistency
- If GC is happening, your „micro“benchmarks need to run longer => Bigger heap, longer running time
- JMH can also be used to construct benchmarks to evaluate (peak) memory consumptions
- Be aware of the different GC implementations and behaviors

Thanks & Enjoy Live!

- JMH extensions available at:
`github.com/cache2k/cache2k-benchmark`
 - `ForcedGcMemoryProfiler`
 - `LinuxVmProfiler`
 - Allocation rate via: `-prof gc`
- Plan/idea: discuss and add it into the JMH code base
- Please like:
`github.com/cache2k/cache2k`

Notes

- How to start visualvm with bigger font size to make screenshots:
`visualvm -J-Dsun.java2d.uiScale=1.0 --fontsize 18`
- First VisualVm heap graph from:
`java-11 -jar jmh-suite/target/benchmarks.jar \\.ZipfianSequenceLoadingBenchmark -jvmArgs -server\ -Xmx10G\ -XX:BiasedLockingStartupDelay=0\ -verbose:gc\ -XX:+UseParallelGC -f 1 -wi -w 15s -i 3 -r 15s -prof comp -prof gc -prof hs_rt -prof hs_gc -prof org.cache2k.benchmark.jmh.LinuxVmProfiler -prof org.cache2k.benchmark.jmh.MiscResultRecorderProfiler -prof org.cache2k.benchmark.jmh.GcProfiler -t 4 -p factor=5 -p entryCount=1000000 -p cacheFactory=org.cache2k.benchmark.Cache2kFactory -rf json -rff /run/shm/jmh-result/result-Cache2kFactory-ZipfianSequenceLoadingBenchmark-4.json`
- Second VisualVm heap graph, additional parameters:
`-prof org.cache2k.benchmark.jmh.ForcedGcMemoryProfiler -gc true`