

# The Missing Benchmark Metric: Memory

2nd February 2019  
Java Dev Room, FOSDEM, Brussels

Jens Wilke  
@cruftex  
cruftex.net

TODO: add copyright!

# About Me



- Performance Fan(atic)
- Author of cache2k <https://cache2k.org>
- 70+ answered questions on StackOverflow about Caching
- JCache / JSR107 Contributor
- General manager of a boutique software engineering shop in Munich
- @cruftex / cruftex.net

# Content

- Why?
- How to gather memory metrics?
- How (not) to gather memory metrics from a running JMH benchmark?
- Real benchmark results ... to get a feeling whether this is worthwhile
- Happyness?!
- Please ask questions right away!

# Motivation

- Compare different caching libraries in identical scenario (access sequence, cache size, CPU and JDK) with JMH
- Primary result: Throughput in **ops/s**
- WT-beep?! Compare **caching libraries** and ignore memory consumption?!
- Space vs. Time trade off
- But:  
no mechanism in JMH to extract memory usage metrics (yet...)

# Metric: Object Graph Traversing

- EHCACHE sizeof
- Java Agent for Memory Measurements

## Pro:

- No actual usage long running benchmark needed, just fill up the data structures
- No complex setup (dedicated hardware)

## Con:

- How to choose the root object / traverse the whole heap?
- Static value of heap objects only

# Metric: Heap Dump / Heap Histogram

- e.g. use jmap to get a heap dump or histogram


## Pro:

- Same as before, we don't need a real benchmark run
- Accurate value of used heap space
- Additional information in the histogram or dump


## Con:

- Costly
- Static value of heap objects only


# Let it Run!



# How to Extract a (single) Metric?








# Side to Side Comparison

Size: 990,380,032 B  
Max: 10,737,418,240 B

Used: 718,396,264 B




^ no forced GC

forced GC >


Size: 1,239,941,120 B  
Max: 10,737,418,240 B

Used: 129,621,464 B



■ Heap size ■ Used heap

# Wait a minute....



# Sidenote:


## GC and Micro Benchmarks

- Micro-Microbenchmark: GC might happen during an iteration
  - Single GC occurrence causes skewed result
  - Use JMH parameter `-gc true` or Zero GC


=> Get garbage collector out of the equation
- NotSo-Microbenchmark: GC happens always during an iteration
  - Make sure a lot of GCs happen, JMH parameter: `-prof gc`
  - Increase warmup and iteration time
  - „Know your GC“  $\leftrightarrow$  „Get to know your GC“

=> You cannot avoid the GC. Make it go steady.

# A Metric? Some Metrics!




# Linux OS Metrics: VmRSS, VmHWM




# GC Notifications

- Notification after GC run (via `GarbageCollectionNotificationInfo`)
- Contains memory usage before and after GC  
(examples only shows notification in iteration 2)




# Allocation Rate





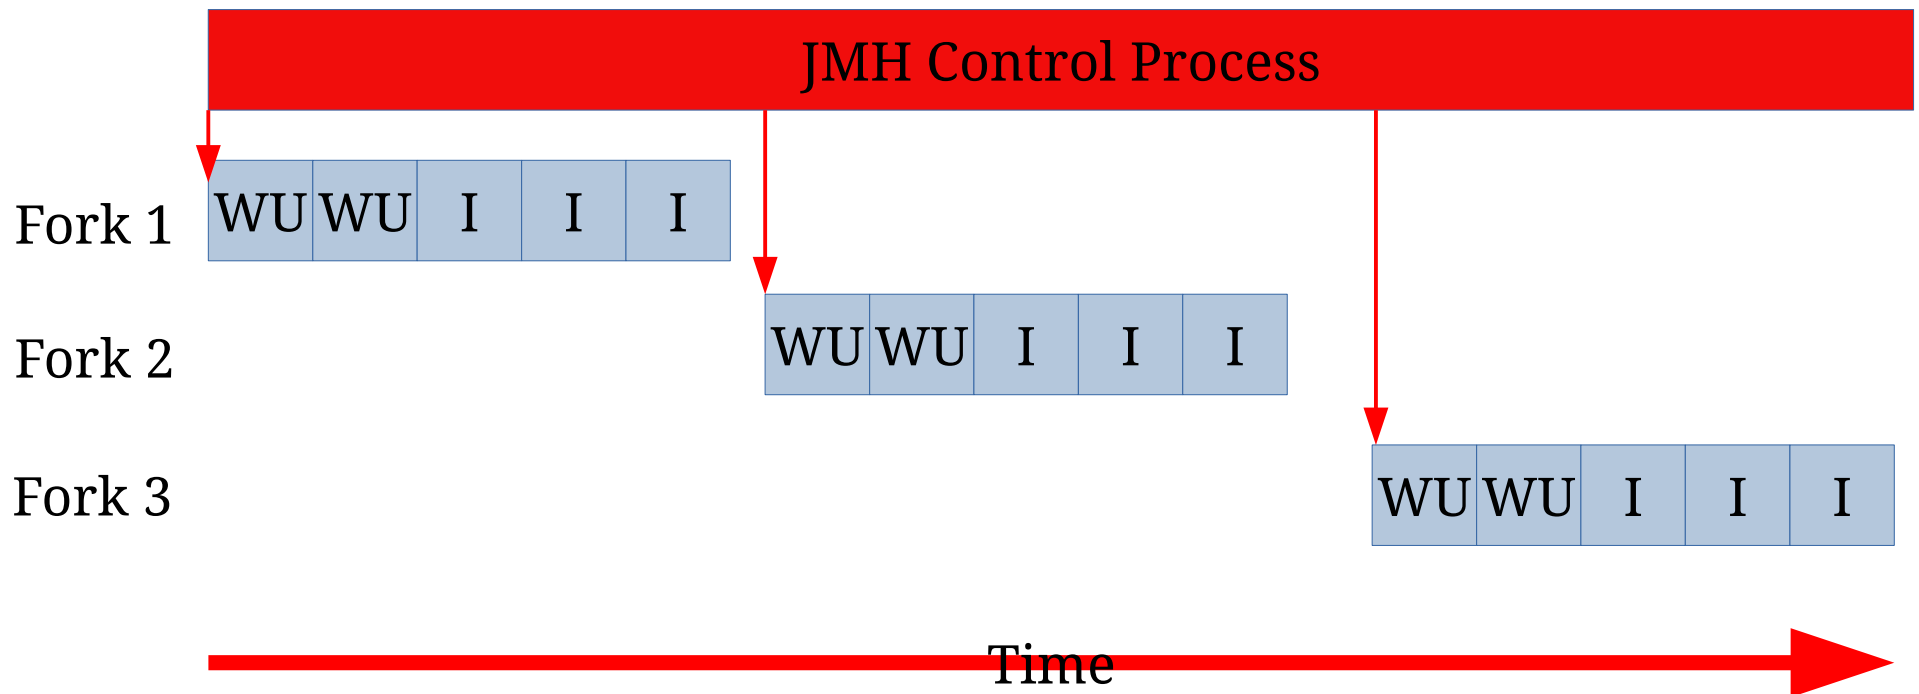
# Minimal Memory at the End

- JMH can run multiple VMs after another (forks)
- at end of each fork the invasive collectors could happen (force GC, heap histogram)



# Example JMH Running Scheme

- JMH can run multiple VMs after another (forks)
- 9 iterations to collect results
- 3 forks, at end of each fork the invasive collectors happen (force GC, heap histogram)
- Parameters: `-f 3 -wi 3 -i 3`



# Results

# JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;

@State(Scope.Thread)
public static class ThreadState {

    ZipfianPattern pattern;

    @Setup(Level.Iteration)
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }
}

@Benchmark @BenchmarkMode(Mode.Throughput)
public long operation(ThreadState threadState, HitCountRecorder rec) {
    // TODO: JMH should return the raw number of operations somewhere...
    rec.opCount++;
    Integer v = cache.get(threadState.pattern.next());
    return v;
}

public Integer load(final Integer key) {
    missCount.increment();
    Blackhole.consumeCPU(1000);
    return key * 2 + 11;
}
```

# JMH Benchmark

```
public int factor = 5;
public int entryCount = 100_000;
```

```
@State(Scope.Thread)
public static class ThreadState {
```

```
    ZipfianPattern pattern;
```

(fast) Zipfian sequence generator per thread  
Skewed random pattern, yields around 90% hitrate

```
    @Setup(Level.Iteration)
```

```
    public void setup(ZipfianSequenceLoadingBenchmark benchmark) {
        pattern = new ZipfianPattern_benchmark.offsetSeed.nextLong(),
            benchmark.entryCount * benchmark.factor);
    }
```

Benchmarked operation is cache.get()  
On miss, cache will invoke a load function(read-trough)

```
@Benchmark @BenchmarkMode({Mode.Throughput})
public long operation(ThreadState threadState, HitCountRecorder rec) {
    // TODO: JMH should return the raw number of operations somewhere...
    rec.opCount++;
    Integer v = cache.get(threadState.pattern.next());
    return v;
}
```

Load function has penalty, via  
Blackhole.consumeCPU()

```
public Integer load(final Integer key) {
    missCount.increment();
    Blackhole.consumeCPU(1000);
    return key * 2 + 11;
}
```

=> Cache eviction is stress for GC!

# Benchmark Setup

## Environment:

- CPU: Intel(R) Xeon(R) CPU E3-1240 v5 @ 3.50GHz, **4 physical cores**
  - Benchmarks run with 4 thread, 4 cores (limited via CPU hotplug)
- Oracle JDK 11, ParallelGC and G1
- Ubuntu 18.04

## Cache Library Versions:

- Google Guava Cache, Version 26
- Caffeine, Version 2.6.2
- cache2k, Version 1.2.0.Final
- EHCache, Version 3.6.1

## Code is at:


<https://github.com/cache2k/cache2k-benchmark>

# JMH Parameters

- JMH Parameters:
  - 3 forks, 2 warmup iterations, 3 measurement iterations, 60 seconds iterations time


=> 9 measurement iterations
- Graphs show the confidence interval
- Confidence interval is at 99.9% confidence level!

# Parallel GC: ops/s







# Parallel GC: heap usage




# Parallel GC: Total / RSS




# G1: heap usage




# G1: Total / RSS



# G1: allocation rate op/s



# G1: allocation rate byte/op



# Conclusion

- The JVM provides different ways to ask for used memory (Runtime, MXBean, jmap / Histogram, GC Notifications)
- If GC is happening, also record memory usage in your JMH results
- JMH can also be used to construct benchmarks to evaluate (peak) memory consumptions
- Be aware of the different GC implementations

# Thanks & Enjoy Live!

- JMH extensions available at:  
`github.com/cache2k/cache2k-benchmark`
  - `ForcedGcMemoryProfiler`
  - `LinuxVmProfiler`
  - Allocation rate via: `-prof gc`
- Plan/idea: discuss and add it into the JMH code base
- Please like:  
`github.com/cache2k/cache2k`



# Notes

- How to start visualvm with bigger font size to make screenshots:  
`visualvm -J-Dsun.java2d.uiScale=1.0 --fontsize 18`
- First VisualVm heap graph from:  
`java-11 -jar jmh-suite/target/benchmarks.jar \\.ZipfianSequenceLoadingBenchmark -jvmArgs -server\ -Xmx10G\ -XX:BiasedLockingStartupDelay=0\ -verbose:gc\ -XX:+UseParallelGC -f 1 -wi -w 15s -i 3 -r 15s -prof comp -prof gc -prof hs_rt -prof hs_gc -prof org.cache2k.benchmark.jmh.LinuxVmProfiler -prof org.cache2k.benchmark.jmh.MiscResultRecorderProfiler -prof org.cache2k.benchmark.jmh.GcProfiler -t 4 -p factor=5 -p entryCount=1000000 -p cacheFactory=org.cache2k.benchmark.Cache2kFactory -rf json -rff /run/shm/jmh-result/result-Cache2kFactory-ZipfianSequenceLoadingBenchmark-4.json`
- Second VisualVm heap graph, additional parameters:  
`-prof org.cache2k.benchmark.jmh.ForcedGcMemoryProfiler -gc true`